

Multiple Linear Regression

The data is located in the file: 'real_estate_price_size_year.csv'.

In this exercise, the dependent variable is 'price', while the independent variables are 'size' and 'year'.

Import the relevant libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sns
sns.set()
```

Load the data

```
In [2]: data = pd.read_csv('real_estate_price_size_year.csv')
```

```
In [3]: data.head()
```

Out[3]:

	price	size	year
0	234314.144	643.09	2015
1	228581.528	656.22	2009
2	281626.336	487.29	2018
3	401255.608	1504.75	2015
4	458674.256	1275.46	2009

```
In [4]: data.describe()
```

Out[4]:

	price	size	year
count	100.000000	100.000000	100.000000
mean	292289.470160	853.024200	2012.600000
std	77051.727525	297.941951	4.729021
min	154282.128000	479.750000	2006.000000
25%	234280.148000	643.330000	2009.000000
50%	280590.716000	696.405000	2015.000000
75%	335723.696000	1029.322500	2018.000000
max	500681.128000	1842.510000	2018.000000

Create the regression

Declare the dependent and the independent variables

```
In [5]: y = data['price']
x1 = data[['size','year']]
```

Regression

```
In [6]: x = sm.add_constant(x1)
results = sm.OLS(y,x).fit()
results.summary()
```

Out[6]:

OLS Regression Results						
Dep. Variable:	price			R-squared:	0.776	
Model:	OLS			Adj. R-squared:	0.772	
Method:	Least Squares			F-statistic:	168.5	
Date:	Wed, 28 Sep 2022			Prob (F-statistic):	2.77e-32	
Time:	18:48:29			Log-Likelihood:	-1191.7	
No. Observations:	100			AIC:	2389.	
Df Residuals:	97			BIC:	2397.	
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-5.772e+06	1.58e+06	-3.647	0.000	-8.91e+06	-2.63e+06
size	227.7009	12.474	18.254	0.000	202.943	252.458
year	2916.7853	785.896	3.711	0.000	1357.000	4476.571
Omnibus:	10.083	Durbin-Watson:	2.250			
Prob(Omnibus):	0.006	Jarque-Bera (JB):	3.678			
Skew:	0.095	Prob(JB):	0.159			
Kurtosis:	2.080	Cond. No.	9.41e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.41e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [ ]:
```